

```

import pandas as pd
import numpy as np
import seaborn as sns
import matplotlib.pyplot as plt
import matplotlib as mpl
from mpl_toolkits.mplot3d import Axes3D
from sklearn.linear_model import LinearRegression
from sklearn.tree import DecisionTreeRegressor
from sklearn.ensemble import RandomForestRegressor
from sklearn.preprocessing import StandardScaler,QuantileTransformer
from sklearn.model_selection import train_test_split
from sklearn.metrics import mean_squared_error, r2_score

%matplotlib inline

traindf = pd.read_csv('train.csv')

traindf.columns

Index(['Id', 'MSSubClass', 'MSZoning', 'LotFrontage', 'LotArea', 'Street',
      'Alley', 'LotShape', 'LandContour', 'Utilities', 'LotConfig',
      'LandSlope', 'Neighborhood', 'Condition1', 'Condition2', 'BldgType',
      'HouseStyle', 'OverallQual', 'OverallCond', 'YearBuilt', 'YearRemodAdd',
      'RoofStyle', 'RoofMatl', 'Exterior1st', 'Exterior2nd', 'MasVnrType',
      'MasVnrArea', 'ExterQual', 'ExterCond', 'Foundation', 'BsmtQual',
      'BsmtCond', 'BsmtExposure', 'BsmtFinType1', 'BsmtFinSF1',
      'BsmtFinType2', 'BsmtFinSF2', 'BsmtUnfSF', 'TotalBsmtSF', 'Heating',
      'HeatingQC', 'CentralAir', 'Electrical', '1stFlrSF', '2ndFlrSF',
      'LowQualFinSF', 'GrLivArea', 'BsmtFullBath', 'BsmtHalfBath', 'FullBath',
      'HalfBath', 'BedroomAbvGr', 'KitchenAbvGr', 'KitchenQual',
      'TotRmsAbvGrd', 'Functional', 'Fireplaces', 'FireplaceQu', 'GarageType',
      'GarageYrBlt', 'GarageFinish', 'GarageCars', 'GarageArea', 'GarageQual',
      'GarageCond', 'PavedDrive', 'WoodDeckSF', 'OpenPorchSF',
      'EnclosedPorch', '3SsnPorch', 'ScreenPorch', 'PoolArea', 'PoolQC',
      'Fence', 'MiscFeature', 'MiscVal', 'MoSold', 'YrSold', 'SaleType',
      'SaleCondition', 'SalePrice', 'TotalBath', 'TotalSF'],
      dtype='object')

numeric_df = traindf.select_dtypes(include='number')
correlation_matrix = numeric_df.corr()
correlation_matrix['SalePrice'].sort_values(ascending = False)

SalePrice      1.000000
OverallQual     0.790982
GrLivArea       0.708624
GarageCars      0.640409
GarageArea      0.623431
TotalBsmtSF     0.613581
1stFlrSF        0.605852
FullBath        0.560664
TotRmsAbvGrd    0.533723
YearBuilt       0.522897
YearRemodAdd     0.507101
GarageYrBlt     0.486362
MasVnrArea      0.477493
Fireplaces      0.466929
BsmtFinSF1      0.386420
LotFrontage     0.351799
WoodDeckSF      0.324413
2ndFlrSF        0.319334
OpenPorchSF     0.315856
HalfBath        0.284108
LotArea         0.263843
BsmtFullBath    0.227122
BsmtUnfSF       0.214479
BedroomAbvGr    0.168213
ScreenPorch     0.111447
PoolArea        0.092404
MoSold          0.046432
3SsnPorch       0.044584
BsmtFinSF2     -0.011378
BsmtHalfBath    -0.016844
MiscVal         -0.021190
Id              -0.021917
LowQualFinSF    -0.025606
YrSold          -0.028923
OverallCond     -0.077856
MSSubClass      -0.084284


```

```
EnclosedPorch    -0.128578
KitchenAbvGr     -0.135907
Name: SalePrice, dtype: float64
```

```
req_tr = ["GarageArea", "OverallQual", "TotalBsmtSF", "1stFlrSF", "2ndFlrSF", "LowQualFinSF", "GrLivArea", "BsmtFullBath", "BsmtHalfBath", "FullBath"]
```

```
selected_tr = traindf[req_tr]
```

```
selected_tr
```




| | GarageArea | OverallQual | TotalBsmtSF | 1stFlrSF | 2ndFlrSF | LowQualFinSF | GrLivArea |
|------|------------|-------------|-------------|----------|----------|--------------|-----------|
| 0 | 548 | 7 | 856 | 856 | 854 | 0 | 1710 |
| 1 | 460 | 6 | 1262 | 1262 | 0 | 0 | 1262 |
| 2 | 608 | 7 | 920 | 920 | 866 | 0 | 1786 |
| 3 | 642 | 7 | 756 | 961 | 756 | 0 | 1717 |
| 4 | 836 | 8 | 1145 | 1145 | 1053 | 0 | 2198 |
| ... | ... | ... | ... | ... | ... | ... | ... |
| 1455 | 460 | 6 | 953 | 953 | 694 | 0 | 1647 |
| 1456 | 500 | 6 | 1542 | 2073 | 0 | 0 | 2073 |
| 1457 | 252 | 7 | 1152 | 1188 | 1152 | 0 | 2340 |
| 1458 | 240 | 5 | 1078 | 1078 | 0 | 0 | 1078 |
| 1459 | 276 | 5 | 1256 | 1256 | 0 | 0 | 1256 |

1460 rows x 15 columns

```
train_df = selected_tr[['TotRmsAbvGrd', 'TotalBath', 'GarageArea', 'TotalSF', 'OverallQual', 'SalePrice']]
```


```
train_df
```



| | TotRmsAbvGrd | TotalBath | GarageArea | TotalSF | OverallQual | SalePrice |
|------|--------------|-----------|------------|---------|-------------|-----------|
| 0 | 8 | 4 | 548 | 4276 | 7 | 208500 |
| 1 | 6 | 3 | 460 | 3786 | 6 | 181500 |
| 2 | 6 | 4 | 608 | 4492 | 7 | 223500 |
| 3 | 7 | 2 | 642 | 4190 | 7 | 140000 |
| 4 | 9 | 4 | 836 | 5541 | 8 | 250000 |
| ... | ... | ... | ... | ... | ... | ... |
| 1455 | 7 | 3 | 460 | 4247 | 6 | 175000 |
| 1456 | 7 | 3 | 500 | 5688 | 6 | 210000 |
| 1457 | 9 | 2 | 252 | 5832 | 7 | 266500 |
| 1458 | 5 | 2 | 240 | 3234 | 5 | 142125 |
| 1459 | 6 | 3 | 276 | 3768 | 5 | 147500 |

1460 rows x 6 columns

```
from sklearn.model_selection import train_test_split
train_set, test_set = train_test_split(train_df, test_size = 0.2, random_state = 42)
print(f"Rows in train set: {len(train_set)}\nRows in test set: {len(test_set)}\n")
```



```
Rows in train set: 1168
Rows in test set: 292
```

```
housing = train_set.drop("SalePrice", axis=1)
housing_labels = train_set["SalePrice"].copy()
```

```
from sklearn.impute import SimpleImputer
from sklearn.pipeline import Pipeline
from sklearn.preprocessing import StandardScaler
my_pipeline = Pipeline([
    ('imputer', SimpleImputer(strategy="median")),
    ('std_scaler', StandardScaler())
])
```

```
X_train = my_pipeline.fit_transform(housing)
X_train
```

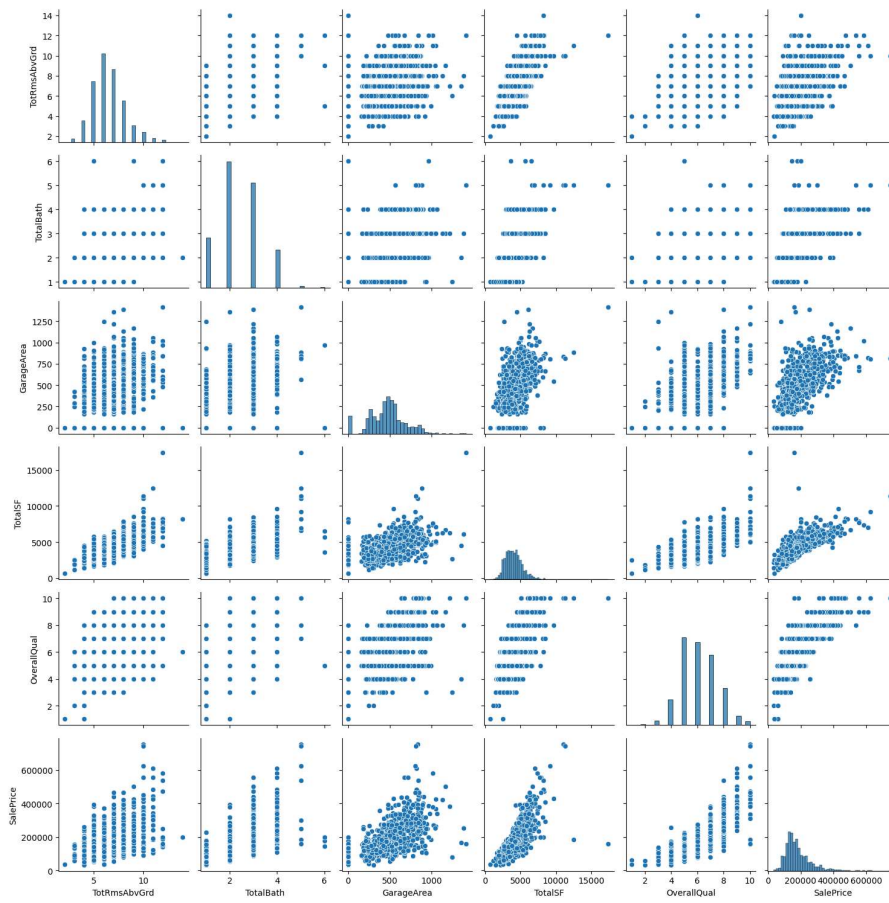
```
-----
NameError                                Traceback (most recent call last)
<ipython-input-3-5243e64da559> in <cell line: 1>()
----> 1 X_train = my_pipeline.fit_transform(housing)
      2 X_train

NameError: name 'my_pipeline' is not defined
```

```
Y_train = housing_labels
Y_train.shape
```

```
(1168,)
```

```
import warnings
warnings.filterwarnings("ignore", category=UserWarning)
%matplotlib inline
sns.pairplot(train_df)
plt.tight_layout()
plt.show()
```




```
corr_matrix = train_df.corr()
corr_matrix['SalePrice'].sort_values(ascending = False)
```

↗

SalePrice1.000000
OverallQual0.790982
TotalSF0.773909
GarageArea0.623431
TotalBath0.613005
TotRmsAbvGrd0.533723
Name: SalePrice, dtype: float64

```
sns.heatmap(train_df.corr(),annot = True)
```



```
testdf = pd.read_csv("test.csv")
testdf.head()
```

↗

| | Id | MSSubClass | MSZoning | LotFrontage | LotArea | Street | Alley | LotShape | LandContou |
|---|------|------------|----------|-------------|---------|--------|-------|----------|------------|
| 0 | 1461 | 20 | RH | 80.0 | 11622 | Pave | NaN | Reg | L\ |
| 1 | 1462 | 20 | RL | 81.0 | 14267 | Pave | NaN | IR1 | L\ |
| 2 | 1463 | 60 | RL | 74.0 | 13830 | Pave | NaN | IR1 | L\ |
| 3 | 1464 | 60 | RL | 78.0 | 9978 | Pave | NaN | IR1 | L\ |
| 4 | 1465 | 120 | RL | 43.0 | 5005 | Pave | NaN | IR1 | HL\ |

5 rows × 10 columns

```
req_tst = ["GarageArea","OverallQual","TotalBsmtSF","1stFlrSF","2ndFlrSF","LowQualFinSF","GrLivArea","BsmtFullBath","BsmtHalfBath","FullBath"
```

```
selected_tst = testdf[req_tst]
```

```
selected_tst.loc[:, 'TotalBath'] = (selected_tst['BsmtFullBath'].fillna(0) + selected_tst['BsmtHalfBath'].fillna(0) + selected_tst['FullBath'
```

```
selected_tst.loc[:, 'TotalSF'] = (selected_tst['TotalBsmtSF'].fillna(0) + selected_tst['1stFlrSF'].fillna(0) + selected_tst['2ndFlrSF'].fill
```


↗

<ipython-input-29-3654aa847672>:1: SettingWithCopyWarning:
A value is trying to be set on a copy of a slice from a DataFrame.
Try using .loc[row_indexer,col_indexer] = value instead

See the caveats in the documentation: https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-c
selected_tst.loc[:, 'TotalBath'] = (selected_tst['BsmtFullBath'].fillna(0) + selected_tst['BsmtHalfBath'].fillna(0) + selected_tst['Fu
<ipython-input-29-3654aa847672>:3: SettingWithCopyWarning:
A value is trying to be set on a copy of a slice from a DataFrame.
Try using .loc[row_indexer,col_indexer] = value instead

See the caveats in the documentation: https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy
 selected_tst.loc[:, 'TotalSF'] = (selected_tst['TotalBsmtSF'].fillna(0) + selected_tst['1stFlrSF'].fillna(0) + selected_tst['2ndFlrSF']


selected_tst



| | GarageArea | OverallQual | TotalBsmtSF | 1stFlrSF | 2ndFlrSF | LowQualFinSF | GrLivArea |
|------|------------|-------------|-------------|----------|----------|--------------|-----------|
| 0 | 730.0 | 5 | 882.0 | 896 | 0 | 0 | 896 |
| 1 | 312.0 | 6 | 1329.0 | 1329 | 0 | 0 | 1329 |
| 2 | 482.0 | 5 | 928.0 | 928 | 701 | 0 | 1629 |
| 3 | 470.0 | 6 | 926.0 | 926 | 678 | 0 | 1604 |
| 4 | 506.0 | 8 | 1280.0 | 1280 | 0 | 0 | 1280 |
| ... | ... | ... | ... | ... | ... | ... | ... |
| 1454 | 0.0 | 4 | 546.0 | 546 | 546 | 0 | 1092 |
| 1455 | 286.0 | 4 | 546.0 | 546 | 546 | 0 | 1092 |
| 1456 | 576.0 | 5 | 1224.0 | 1224 | 0 | 0 | 1224 |
| 1457 | 0.0 | 5 | 912.0 | 970 | 0 | 0 | 970 |
| 1458 | 650.0 | 7 | 996.0 | 996 | 1004 | 0 | 2000 |

1459 rows × 14 columns

```
test_df_unproc = selected_tst[['TotRmsAbvGrd', 'TotalBath', 'GarageArea', 'TotalSF', 'OverallQual']]
test_df_unproc
```




| | TotRmsAbvGrd | TotalBath | GarageArea | TotalSF | OverallQual |
|------|--------------|-----------|------------|---------|-------------|
| 0 | 5 | 1.0 | 730.0 | 2674.0 | 5 |
| 1 | 6 | 2.0 | 312.0 | 3987.0 | 6 |
| 2 | 6 | 3.0 | 482.0 | 4186.0 | 5 |
| 3 | 7 | 3.0 | 470.0 | 4134.0 | 6 |
| 4 | 5 | 2.0 | 506.0 | 3840.0 | 8 |
| ... | ... | ... | ... | ... | ... |
| 1454 | 5 | 2.0 | 0.0 | 2730.0 | 4 |
| 1455 | 6 | 2.0 | 286.0 | 2730.0 | 4 |
| 1456 | 7 | 2.0 | 576.0 | 3672.0 | 5 |
| 1457 | 6 | 2.0 | 0.0 | 2852.0 | 5 |
| 1458 | 9 | 3.0 | 650.0 | 4996.0 | 7 |

1459 rows × 5 columns

```
test_df = test_df_unproc.fillna(test_df_unproc.mean())
```

```
x_test = my_pipeline.transform(test_df[['TotRmsAbvGrd', 'TotalBath', 'GarageArea', 'TotalSF', 'OverallQual']].values)
x_test
```



```
array([[ -0.96456591, -1.57881784,  1.2024646 , -1.10333489, -0.82044456],
       [ -0.34690528, -0.48377079, -0.77853123, -0.09910341, -0.08893368],
       [ -0.34690528,  0.61127627,  0.02713693,  0.05309923, -0.82044456],
       ...,
       [  0.27075534, -0.48377079,  0.47262403, -0.34002719, -0.82044456],
       [ -0.34690528, -0.48377079, -2.25716927, -0.96719384, -0.82044456],
       [  1.50607659,  0.61127627,  0.82332664,  0.67261751,  0.64257719]])
```

```
#model = LinearRegression()
#model = DecisionTreeRegressor()
model = RandomForestRegressor()
model.fit(X_train, Y_train)
```

```

RandomForestRegressor
RandomForestRegressor()

```

```
y_train_pred = model.predict(X_train)
```

```
y_train_pred[:5]
```

```
array([148521.18, 172056.9 , 90154. , 166354.87, 136974. ])
```

```
some_data = housing.iloc[:5]
```

```
some_labels = housing_labels.iloc[:5]
```

```
proc_data = my_pipeline.transform(some_data)
```

```
model.predict(proc_data)
```

```
array([148521.18, 172056.9 , 90154. , 166354.87, 136974. ])
```

```
list(some_labels)
```

```
[145000, 178000, 85000, 175000, 127000]
```

```
train_mse = mean_squared_error(Y_train,y_train_pred)
```

```
train_rmse = np.sqrt(train_mse)
```

```
print(f"Training MSE: {train_mse:.2f}, Training RMSE: {train_rmse:.2f}")
```

```
Training MSE: 163787325.95, Training RMSE: 12797.94
```

```
from sklearn.model_selection import cross_val_score
```

```
scores = cross_val_score(model,X_train,Y_train,scoring="neg_mean_squared_error",cv = 200)
```

```
rmse_scores = np.sqrt(-scores)
```

```
rmse_scores
```

```

array([ 20134.98781949, 14686.87903424, 25562.67014137, 11905.34310038,
        45553.74485603, 12497.36432866, 20265.04862479, 12889.92955452,
        11261.55705183, 52747.54088485, 34609.31832455, 29043.5516443 ,
        14251.0653049 , 9526.57305502, 20334.1112828 , 21620.68600092,
        19603.56284842, 32427.16875131, 36371.87422057, 22903.17476099,
        29016.6143909 , 17732.6279336 , 17531.1583017 , 27387.69596271,
        19728.8207265 , 19055.45171987, 42351.70427429, 38289.73939472,
        169442.94448404, 51948.6841921 , 20039.87728897, 31810.0452124 ,
        19298.84508228, 31237.14614987, 48310.76375112, 12092.20975801,
        24145.4982442 , 29326.85866748, 19506.94453812, 31756.81340151,
        25479.30428831, 32190.70892996, 29135.34518748, 33369.33615978,
        33584.70848518, 31432.59335172, 26504.67187058, 43618.47998814,
        20860.76813753, 20657.5541872 , 21071.80558272, 53215.5982509 ,
        38228.71510965, 36641.28538325, 22073.26195159, 29934.96832478,
        9650.5560506 , 25607.4089106 , 22936.90169392, 30883.30469553,
        196725.30642846, 13491.25136136, 24715.70765939, 35421.61088696,
        23144.38847077, 24795.38663266, 41835.06611542, 30494.15015245,
        6555.91515732, 15579.91845922, 56876.132711 , 25937.21951802,
        57422.75017867, 19575.1187025 , 46137.92518233, 45489.93083409,
        37558.36328828, 28587.41603817, 38103.47081189, 32716.94616841,
        26202.3327639 , 32227.48042114, 107726.2272588 , 10285.87273477,
        48266.92082725, 31804.01462906, 18613.94256777, 23131.08842726,
        36113.34216564, 78378.17033651, 18775.59369568, 21453.69985058,
        25668.55756481, 24129.84654464, 21071.21765313, 21733.66626709,
        24587.97829035, 26258.57635744, 68989.42232601, 12352.74900534,
        32906.76054716, 27480.11429687, 45547.59333722, 54560.41029764,
        18115.99299691, 19096.94691454, 36792.92575559, 18896.67704669,
        19276.91598619, 14519.23555926, 19970.06860648, 12423.77754214,
        12094.27428157, 18969.35804675, 29244.15652398, 33901.53161987,
        75631.53392333, 28047.73904368, 18379.3750864 , 26041.29291982,
        28324.9575722 , 29756.48018735, 14313.75217015, 18665.93479935,
        31275.61234144, 21641.25100101, 28743.87731493, 46061.1868786 ,
        37400.46496854, 15846.09278188, 36268.70180244, 7769.55188641,
        25458.80855174, 26958.30089985, 28917.79069789, 11310.055774 ,
        47552.44402414, 33432.36993426, 30679.33718408, 12573.17246057,

```

```

24470.12191997, 24497.22777048, 26025.19328146, 14213.07304743,
20817.14421494, 29400.88794654, 32918.50467223, 12146.96259758,
12853.61488526, 14973.26035789, 23699.98583852, 25934.69156306,
13975.19548295, 42630.04201458, 29441.60406063, 13362.67776993,
26307.66408174, 25364.09111996, 20196.22774016, 25136.92955551,
57581.48168316, 27483.57960713, 26304.55335392, 16805.65646023,
48771.61740836, 18986.33355552, 12223.46789884, 61086.90902796,
21223.61539812, 30941.98171267, 31486.9433332, 22021.52735504,
15835.77256873, 12102.93082367, 20000.9617597, 15518.48566011,
22174.64973984, 15999.94009796, 35208.88688937, 21299.11173005,
36057.75718532, 10168.83999179, 16645.47265271, 7039.58608775,
19899.05658098, 28872.26749384, 26077.70350845, 12630.91552385,
26073.00929526, 41985.17469325, 29530.10815622, 19655.47974547,
17120.754699, 12897.45423417, 17405.99170086, 17088.11487907,
59854.26733955, 20239.14386941, 25271.46448594, 29172.88345019])

```

```

def print_scores(scores):
    print("Scores:", scores)
    print("Mean:", scores.mean())
    print("Standard Deviation", scores.std())

```

```
print_scores(rmse_scores)
```

```

Scores: [ 20134.98781949 14686.87903424 25562.67014137 11905.34310038
45553.74485603 12497.36432866 20265.04862479 12889.92955452
11261.55705183 52747.54088485 34609.31832455 29043.5516443
14251.0653049 9526.57305502 20334.1112828 21620.68600092
19603.56284842 32427.16875131 36371.87422057 22903.17476099
29016.6143909 17732.6279336 17531.1583017 27387.69596271
19728.8207265 19055.45171987 42351.70427429 38289.73939472
169442.94448404 51948.6841921 20039.87728897 31810.0452124
19298.84508228 31237.14614987 48310.76375112 12092.20975801
24145.4982442 29326.85866748 19506.94453812 31756.81340151
25479.30428831 32190.70892996 29135.34518748 33369.33615978
33584.70848518 31432.59335172 26504.67187058 43618.47998814
20860.76813753 20657.5541872 21071.80558272 53215.5982509
38228.71510965 36641.28538325 22073.26195159 29934.96832478
9650.5560506 25607.4089106 22936.90169392 30883.30469553
196725.30642846 13491.25136136 24715.70765939 35421.61088696
23144.38847077 24795.38663266 41835.06611542 30494.15015245
6555.91515732 15579.91845922 56876.132711 25937.21951802
57422.75017867 19575.1187025 46137.92518233 45489.93083409
37558.36328828 28587.41603817 38103.47081189 32716.94616841
26202.3327639 32227.48042114 107726.2272588 10285.87273477
48266.92082725 31804.01462906 18613.94256777 23131.08842726
36113.34216564 78378.17033651 18775.59369568 21453.69985058
25668.55756481 24129.84654464 21071.21765313 21733.66626709
24587.97829035 26258.57635744 68989.42232601 12352.74900534
32906.76054716 27480.11429687 45547.59333722 54560.41029764
18115.99299691 19096.94691454 36792.92575559 18896.67704669
19276.91598619 14519.23555926 19970.06860648 12423.77754214
12094.27428157 18969.35804675 29244.15652398 33901.53161987
75631.53392333 28047.73904368 18379.3750864 26041.29291982
28324.9575722 29756.48018735 14313.75217015 18665.93479935
31275.61234144 21641.25100101 28743.87731493 46061.1868786
37400.46496854 15846.09278188 36268.70180244 7769.55188641
25458.80855174 26958.30089985 28917.79069789 11310.055774
47552.44402414 33432.36993426 30679.33718408 12573.17246057
24470.12191997 24497.22777048 26025.19328146 14213.07304743
20817.14421494 29400.88794654 32918.50467223 12146.96259758
12853.61488526 14973.26035789 23699.98583852 25934.69156306
13975.19548295 42630.04201458 29441.60406063 13362.67776993
26307.66408174 25364.09111996 20196.22774016 25136.92955551
57581.48168316 27483.57960713 26304.55335392 16805.65646023
48771.61740836 18986.33355552 12223.46789884 61086.90902796
21223.61539812 30941.98171267 31486.9433332 22021.52735504
15835.77256873 12102.93082367 20000.9617597 15518.48566011
22174.64973984 15999.94009796 35208.88688937 21299.11173005
36057.75718532 10168.83999179 16645.47265271 7039.58608775
19899.05658098 28872.26749384 26077.70350845 12630.91552385
26073.00929526 41985.17469325 29530.10815622 19655.47974547
17120.754699 12897.45423417 17405.99170086 17088.11487907
59854.26733955 20239.14386941 25271.46448594 29172.88345019]
Mean: 28970.871605843375
Standard Deviation 20827.428551337565

```

```
y_pred=model.predict(x_test)
```

```
y_pred
```



```
array([130460.83, 155886.5, 145079. , ..., 138131.5, 110141.5 ,
       235697.2 ])
```

```
pred=pd.DataFrame(y_pred)
sub_df=pd.read_csv('sample_submission.csv')
datasets=pd.concat([sub_df['Id'],pred],axis=1)
datasets.columns=['Id','SalePrice']
datasets.to_csv('sample_submission.csv',index=False)
```

```
traindf['TotalBath'] = traindf['FullBath'] + traindf['HalfBath'] + traindf['BsmtFullBath'] + traindf['BsmtHalfBath']
traindf['TotalSF'] = traindf['1stFlrSF'] + traindf['2ndFlrSF'] + traindf['TotalBsmtSF']
```

```
req_tr = ["GarageArea", "OverallQual", "TotalBath", "TotalSF", "TotRmsAbvGrd", "SalePrice"]
selected_tr = traindf[req_tr].fillna(0)
```

```
train_set, test_set = train_test_split(selected_tr, test_size=0.2, random_state=42)
```

```
X_train = train_set.drop("SalePrice", axis=1)
Y_train = train_set["SalePrice"].copy()
```

```
X_test = test_set.drop("SalePrice", axis=1)
Y_test = test_set["SalePrice"].copy()
```

```
pipeline = Pipeline([
    ('imputer', SimpleImputer(strategy="median")),
    ('std_scaler', StandardScaler())
])
```

```
X_train_prepared = pipeline.fit_transform(X_train)
X_test_prepared = pipeline.transform(X_test)
```

```
model = LinearRegression()
model.fit(X_train_prepared, Y_train)
```

```
LinearRegression()
```


```
Y_train_pred = model.predict(X_train_prepared)
train_mse = mean_squared_error(Y_train, Y_train_pred)
train_rmse = np.sqrt(train_mse)
train_r2 = r2_score(Y_train, Y_train_pred)
print(f"Training MSE: {train_mse:.2f}, Training RMSE: {train_rmse:.2f}, Training R^2: {train_r2:.2f}")
```

```
Training MSE: 1451670460.71, Training RMSE: 38100.79, Training R^2: 0.76
```

```
Y_test_pred = model.predict(X_test_prepared)
test_mse = mean_squared_error(Y_test, Y_test_pred)
test_rmse = np.sqrt(test_mse)
test_r2 = r2_score(Y_test, Y_test_pred)
print(f"Test MSE: {test_mse:.2f}, Test RMSE: {test_rmse:.2f}, Test R^2: {test_r2:.2f}")
```

```
Test MSE: 1568951446.76, Test RMSE: 39609.99, Test R^2: 0.80
```

```
plt.figure(figsize=(10, 6))
plt.scatter(Y_test, Y_test_pred, color='blue', alpha=0.6)
plt.title('Actual vs Predicted House Prices')
plt.xlabel('Actual Price')
plt.ylabel('Predicted Price')
plt.xlim(0, max(Y_test.max(), Y_test_pred.max()))
plt.ylim(0, max(Y_test.max(), Y_test_pred.max()))
plt.plot([0, max(Y_test.max(), Y_test_pred.max())], [0, max(Y_test.max(), Y_test_pred.max())], color='red', linestyle='--')
plt.grid(True)
plt.show()
```



```
errors = Y_test_pred - Y_test
plt.figure(figsize=(10, 6))
sns.histplot(errors, bins=30, kde=True, color='blue')
plt.title('Distribution of Prediction Errors')
plt.xlabel('Prediction Error')
plt.ylabel('Count')
plt.xlim(-100000, 100000)
```